

Creating process in UNIX/LINUX

- Use `fork()` system call to create new processes
- The process making the system call is the parent
- The new process is called the child process
- Child process receives a copy of the address space of the parent
- The two processes are distinguished by the value returned
- A value of 0 is returned to the child and the ID of the child process is returned to the parent
- Both processes continue execution at the instruction after the `fork()`

Example 1

```
#include<stdio.h>
#include<sys/types.h>
#include<unistd.h>
void main(){
    int  pid;
    pid = fork();
    switch (pid) {
    case -1 : printf(" fork failed");
        exit(-1);
    case 0 : printf("I am the child, ID = %d\n", getpid());
        exit(0);
    default : printf("I am the parent, ID = %d\n", getpid());
    }
}
```

Example 2

```
#include<stdio.h>
#include<sys/types.h>
#include<unistd.h>
void main(){
    int  i, n;
    pid_t  pid;
    for (i = 0; i < n; i++)
        if (pid = fork() == 0)
            break;
    printf("My process ID = %d\n and my parent's ID =
%d\n",  getpid(),  getppid());
}
```

- This can be used to create n children of a parent

wait() System Call

- The process making the wait() system call waits until its child completes or stops or until the caller receives a signal
- The *wait* returns immediately if the process has no child
- If the child terminates, the value returned is id of the child and is greater than 0
- wait() takes one argument which is an integer pointer and it stores the return status of the child

Example

```
#include <stdio.h>
#include <sys/wait.h>
#include <sys/types.h>
#include <unistd.h>
void main(){
    pid_t pid;
    int status;
    pid = fork();
    switch (pid) {
        case -1 : printf(" fork failed");
            exit(-1);
        case 0 : printf("I am the child, ID = %d\n", getpid());
            exit(0);
        default : if (wait (&status) != pid)
            printf("A signal has interrupted the wait);
            else
            printf("I am the parent, ID = %d and child ID =%d\n", getpid());
    }
}
```

exec() System Call

- fork() creates a duplicate of the calling process
- May require child to work on a different program
- exec() system call allows a new process to take the place of the calling process
- The new process starts executing at its main() function
- Usually use a combination of fork()-exec() sequence to let child execute a different program while the parent continues with the original program

Variations of exec()

- Six different functions
 - Differentiates the way command line arguments are given and the name of the program file is specified
- The six functions are:
 - execl()
 - List arguments with the call; list ends with (char *)0
 - execv()
 - Place arguments in a vector

Variations of exec()

- `execle()`
 - Same as `execl()` with additional argument that specifies the environment of the new program. In `exec()` forms in which environment is not provided the new program takes the environment of the parent
- `execvp()`
 - Same as `execl()`, but uses the `PATH` environment variable to search for the executable.

Variations of exec()

- `execvp()`
 - Same as `execv()`; uses the `PATH` environment variable to search for the executable
- `execve()`
 - Same as `execv()`; takes additional variable for the environment of the new program

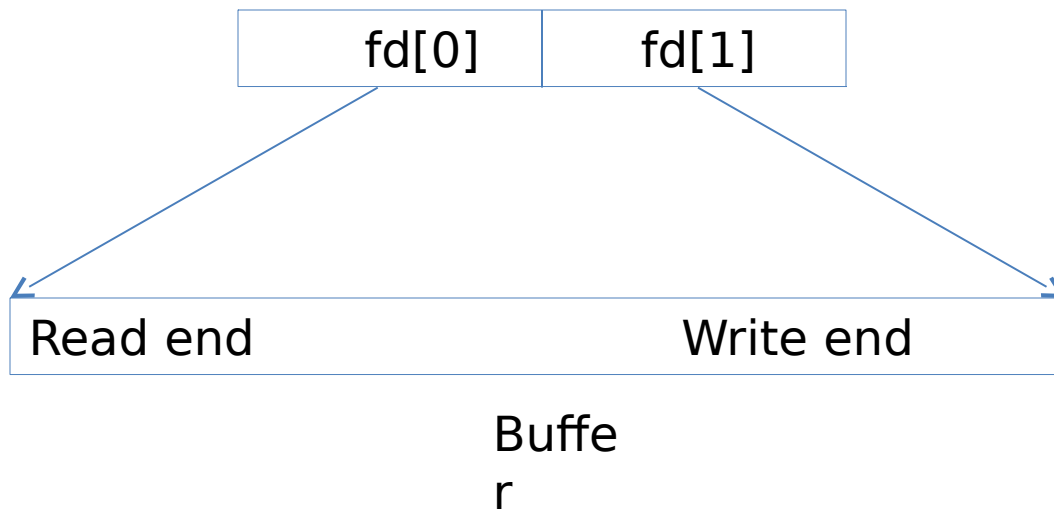
Example

```
#include    <stdio.h>
#include    <sys/types.h>
#include    <unistd.h>
#include    <stdlib.h>
#include    <sys/wait.h>
void main(){
    pid_t pid;
    int  status;
    pid = fork();
    switch (pid) {
    case -1 : printf(" fork failed");
              exit(-1);
    case 0 :  printf("I am the child, ID = %d\n", getpid());
              execl("/usr/bin/ls", "ls", -l, NULL);
    default : if (wait (&status) != pid)
                printf("A signal has interrupted the wait);
                else
                printf("I am the parent, ID = %d and child ID =%d\n", getpid());
    }
}
```

pipe() System Call

- Used for communications between parent and child
- The pipe(fd) system call creates a communication buffer, where fd[] is an integer array of two elements
- It returns two file descriptors fd[0] and fd[1].
- To read from the file, use fd[0]
- To write to the file, use fd[1]
- Use read() system call to read from the buffer
- Use write() system call to write to the buffer
- Buffer is organized as a FIFO queue

pipe() system call buffer



read()/write() system calls

- read() and write() system calls are defined as follows:

`ssize_t read(int fd, void *buf, size_t nbyte)`

read() requests to read nbyte bytes from the file pointed to by fd and placed in buf. buf should be large enough to hold nbyte bytes; returns number of bytes actually read

`ssize_t write(int fd, const void *buf, size_t nbyte)`

write() requests to write nbyte bytes from buf to the file with descriptor fd; returns the number of bytes actually written

Example

- `#include <errno.h>`
- `#include <stdio.h>`
- `#include <unistd.h>`
- `#include <stdlib.h>`
- `#include <sys/types.h>`
- `#define bufsize 20`
- `int fildes[2];`
- `char buffer[bufsize];`
- `void error(void)`
- `{`
- `printf("unable to create a new process - job terminated \n");`
- `fflush(stdout);`
- `}`

Example contd.

```
void child(void)
{
    int pm, i;
    printf("child starts working.\n");
    close(fildes[1]);
    for (i = 0; i<5; i++)
    {
        printf("in child i = %d \n", i);
        pm = read(fildes[0],&buffer,20);
        printf ("child received %d chars from parent\n",pm);
        printf ("message received is: %s ",buffer);
        fflush(stdout);
        sleep(1);
    }
    exit(0);
}
```

Example contd.

```
void parent(void)
{
    int i, pm;
    printf("parent starts working.\n");
    close(fildes[0]);
    for (i = 0; i < 5 ; i++)
    {
        switch(i) {
            case 0:
                pm = write(fildes[1], "first message \n", 20);
                printf("characters written by parent = %d \n", pm);
                fflush(stdout);
                break;
```


Example contd.

case 1:

```
pm = write(fildes[1], "second message \n", 20);  
printf("characters written by parent = %d \n", pm);  
fflush(stdout);
```

break;

case 2:

```
pm = write(fildes[1], "third message \n", 20);  
printf("characters written by parent = %d \n", pm);  
fflush(stdout);
```

break;

Example contd.

case 3:

```
pm = write(fildes[1], "fourth message \n", 20);  
printf("characters written by parent = %d \n", pm);  
fflush(stdout);
```

break;

case 4:

```
pm = write(fildes[1], "fifth message \n", 20);  
printf("characters written by parent = %d \n", pm);  
fflush(stdout);
```

```
}
```

```
}
```

```
wait(NULL);  
printf("parent terminating\n");  
exit(0);
```

```
}
```

Example contd.

```
int main(void)
{
    int pp, pid;

    pp = pipe (fildes);
    printf("pipe opened with pp = %d \n", pp);
    fflush(stdout);
    pid = fork();
    if (pid == 0)
        child();
    else if (pid == -1)
        error();
    else parent();
}
```

Pipe() System call

- pipe() creates a unidirectional communication buffer
- To have a two-way communication, need to make two pipe() system calls
- For read(), end-of-file condition for pipes occurs only when pipe is empty and there are no more processes with write descriptors open to the pipe. Need to close all descriptors not needed
- If a pipe is empty and write descriptors to the pipe are open, the process will wait